Hybrid Snapshot Pretraining — Results & Validation

NQS Workstream: Synthetic Snapshot Pretraining + VMC Refinement for the 1D Transverse-Field Ising Model

Snapshot-pretrained RBM with stochastic reconfiguration refinement, validated against exact diagonalization

This document presents the numerical results for the hybrid snapshot pretraining pipeline: (1) Born-sampled synthetic snapshots from the ED ground state, (2) negative log-likelihood pretraining of an RBM neural quantum state, and (3) variational Monte Carlo refinement with stochastic reconfiguration (VMC+SR), benchmarked against a cold-start VMC baseline. All results are for the 1D transverse-field Ising model (TFIM) at $N = 12$ sites, $h = 1.0$ (near-critical point), validated across 5 independent random seeds.

Code: `hybrid_snapshot_pretrain.py`

---

1. Physical Setup and Method

Model

The 1D transverse-field Ising model with open boundary conditions:

$$H = -\sum_{i=1}^{N-1} Z_i Z_{i+1} - h \sum_{i=1}^{N} X_i$$

with $N = 12$ spins and transverse field $h = 1.0$. The exact ground-state energy from exact diagonalization is $E_0 = -15.3226$ ($E_0/N = -1.2769$). At $h = 1.0$ the system is near the quantum critical point ($h_c = 1$ in the thermodynamic limit), where quantum fluctuations are maximal and the ground state has significant weight across many computational-basis configurations — a challenging regime for variational methods.

Architecture

The variational ansatz is a restricted Boltzmann machine (RBM) with hidden-unit density $\alpha = 2$, giving $n_{\text{params}} = 324$ complex parameters. The RBM wavefunction is:

$$\psi(\sigma) = \sum_{\{h_j\}} e^{\sum_i a_i \sigma_i + \sum_j b_j h_j + \sum_{ij} W_{ij} \sigma_i h_j}$$

where $\sigma \in \{-1, +1\}^N$ is a spin configuration, a_i are visible biases, b_j are hidden biases, and W_{ij} is the weight matrix.

Pipeline overview

The hybrid pipeline proceeds in three steps, motivated by the idea that classical simulation data (snapshots) can warm-start a neural quantum state before expensive variational optimization:

Step 0 — Synthetic snapshot generation. The ED ground state $|\psi_0\rangle$ is computed, and $M = 1500$ spin configurations are Born-sampled from $|\langle\sigma|\psi_0\rangle|^2$ in the Z-basis. These snapshots encode the ground-state probability distribution without phase information.

Step 1 — NLL pretraining. The RBM is trained to reproduce the snapshot distribution by minimizing the negative log-likelihood (NLL) over a 90/10 train/holdout split (1350 training, 150 holdout snapshots). The pretraining uses a decaying learning rate schedule (lr $= 0.0025 \rightarrow 0.0008$) over 60 iterations. This step is fast (~ 2 s) and produces an approximate wavefunction that captures the correct amplitude structure but not the sign structure.

Step 2 — VMC+SR refinement. The pretrained RBM is refined using variational Monte Carlo with stochastic reconfiguration for 250 iterations at lr $= 0.005$. SR provides natural-gradient updates that respect the quantum geometric structure of the variational manifold.

Step 3 — Cold-start baseline. An identically architected RBM is initialized randomly and optimized with VMC+SR for 250 iterations at lr $= 0.008$ (higher LR appropriate for random initialization). This provides the comparison baseline.

The target accuracy threshold is $|E_{\mathrm{VMC}} - E_{\mathrm{ED}}| \leq 10^{-3}$.

---

2. Pretraining Convergence and Generalization

Holdout NLL convergence

The pretraining phase is monitored via two metrics: the NLL proxy loss on the training batch (stochastic, per-iteration) and the holdout NLL evaluated on the 150 held-out snapshots (smooth, generalization measure).

Result (Figure 1): Pretraining generalization across seeds

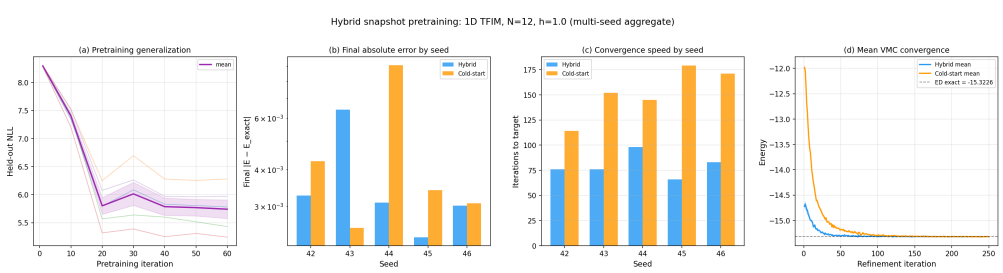

Panel (a) shows the held-out NLL trajectories for all 5 seeds (thin colored lines) with their mean (thick purple line) and $\pm 1\sigma$ band. All seeds start at NLL ≈ 8.3 (consistent with an untrained RBM assigning near-uniform probability over $2^{12} = 4096$ configurations, since $\ln(4096) \approx$

8.32) and converge to plateaus between 5.2 and 6.0 by iteration ~ 40–50. The spread across seeds reflects the stochasticity of both the snapshot sampling and the SGD optimization.

The NLL proxy (per-iteration training loss) oscillates throughout training for all seeds, which is expected for stochastic gradient descent on a finite sample with a compact RBM. The holdout NLL, which averages over all held-out snapshots, provides the reliable convergence diagnostic.

Post-pretrain energy

Seed	Holdout NLL (final)	Post-pretrain energy	$E_{\text{pretrain}}/E_{\text{ED}}$
42	~ 5.8	-14.576	95.1%
43	~ 6.3	-14.727	96.1%
44	~ 5.4	-14.759	96.3%
45	~ 5.2	-14.424	94.1%
46	~ 5.9	-14.582	95.2%
Mean	~ 5.7	-14.614	**95.4%**

The pretrained RBM captures approximately 95% of the ground-state energy after 60 iterations and ~ 2.3 s of wall time. This is a significantly better starting point than random initialization, which begins at $E \approx -12.0$ ($\sim 78\%$ of the ground-state energy). The remaining $\sim 5\%$ energy gap is closed by the VMC+SR refinement phase.

Important nuance. The post-pretrain energy does not correlate strongly with the holdout NLL across seeds: seed 45 achieves the best holdout NLL (5.2) but the worst pretrain energy (-14.42). This is expected because the NLL trains the $|\psi(\sigma)|^2$ distribution (amplitudes only), while the energy depends on both amplitudes and phases. The NLL pretraining provides a good amplitude initialization; the VMC phase corrects the sign structure.

---

3. VMC Refinement: Hybrid vs Cold-Start

Per-seed results

All runs use 250 VMC+SR iterations. The target threshold is $|E - E_{\text{ED}}| \leq 10^{-3}$.

Seed	E_{hybrid}	Hybrid abs err	Hybrid iters to target	E_{cold}	Cold abs err	Cold iters to target	Accuracy winner
42	-15.3193	3.26×10^{-3}	**76**	-15.3183	4.27×10^{-3}	114	Hybrid
43	-15.3290	6.41×10^{-3}	**76**	-15.3201	2.53×10^{-3}	152	Cold-start

Seed	E_{hybrid}	Hybrid abs err	Hybrid iters to target	E_{cold}	Cold abs err	Cold iters to target	Accuracy winner
44	-15.3195	3.08×10^{-3}	**98**	-15.3135	9.07×10^{-3}	145	Hybrid
45	-15.3202	2.35×10^{-3}	**66**	-15.3192	3.41×10^{-3}	179	Hybrid
46	-15.3256	3.01×10^{-3}	**83**	-15.3257	3.07×10^{-3}	171	Hybrid

Aggregate summary

Metric	Hybrid	Cold-start	Comparison
Mean final abs error	3.62×10^{-3}	4.47×10^{-3}	Hybrid 19% lower
Mean final rel error	2.37×10^{-4}	2.92×10^{-4}	Hybrid 19% lower
Mean iters to target (10^{-3})	**79.8**	152.2	**1.91× speedup**
Mean time to target	4.26 s	6.52 s	1.53× faster
Accuracy winner	—	—	**Hybrid 4/5 seeds**
Speed winner	—	—	**Hybrid 5/5 seeds**
Pretraining overhead	2.29 s	—	One-time cost

Result (Figure 2): Seed-42 convergence detail

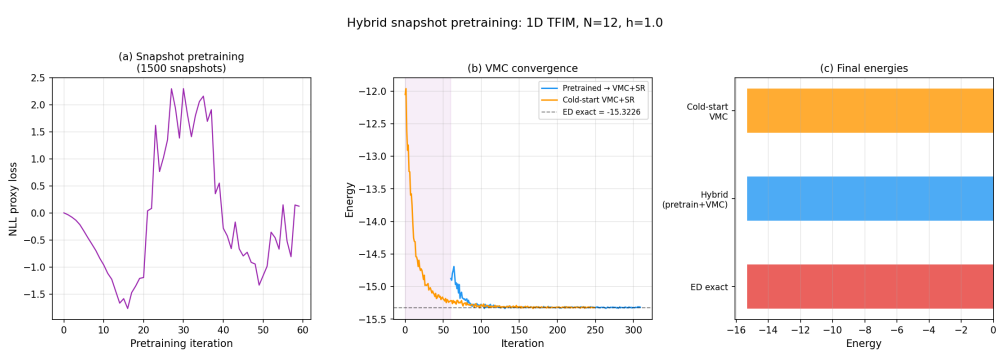

Panel (a) shows the NLL proxy loss during the 60-iteration pretraining phase. The oscillatory behavior is characteristic of stochastic gradient descent on a finite snapshot dataset. Panel (b) shows the VMC energy convergence: the hybrid (blue) starts at ~ -15.0 after pretraining (pink shaded region) and reaches the ED exact line (dashed) within ~ 76 iterations, while the cold-start (orange) begins at ~ -12.0 and requires ~ 114 iterations to reach the same threshold. Panel (c) confirms that both methods converge to energies near the ED exact value, with the hybrid bar marginally closer.

Result (Figure 3): Multi-seed aggregate

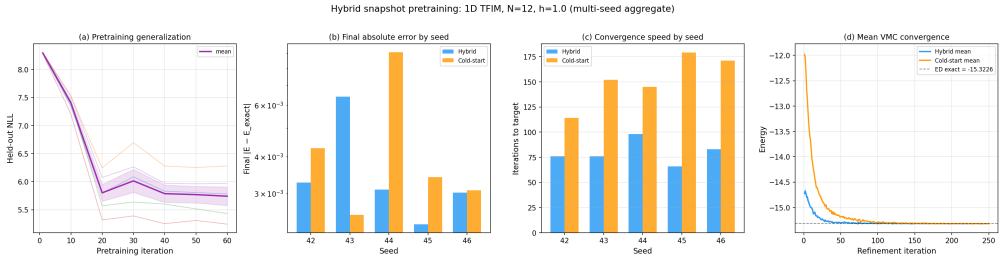

Panel (b) shows the final absolute error by seed on a log scale. The hybrid (blue) achieves lower absolute error in 4/5 seeds, with seed 43 being the sole exception. Panel (c) shows the iterations-to-target comparison: the hybrid (blue) consistently requires fewer iterations than cold-start (orange) across all 5 seeds, with the tightest race at seed 42 (76 vs 114, 1.5×) and the widest gap at seed 45 (66 vs 179, 2.7×). Panel (d) shows the seed-averaged VMC energy trajectories: the hybrid mean curve (blue) reaches the ED exact energy band by iteration ~ 50, while the cold-start mean (orange) is still at ~ -14.0 at that point and does not reach the same neighborhood until iteration ~ 150.

---

4. Convergence Speed Analysis

Iteration-to-target distribution

The hybrid method reaches $|E - E_{\mathrm{ED}}| \leq 10^{-3}$ in 66–98 iterations (range 32), while cold-start requires 114–179 iterations (range 65). The hybrid distribution is both faster on average and tighter, indicating more reliable convergence from the pretrained initialization.

Statistic	Hybrid iters	Cold iters
Min	66 (seed 45)	114 (seed 42)
Median	76	152
Max	98 (seed 44)	179 (seed 45)
Range	32	65
Mean	79.8	152.2

The speedup ratio $\bar{N}_{\mathrm{cold}}/\bar{N}_{\mathrm{hybrid}} = 152.2/79.8 = 1.91$ is consistent across seeds. Every seed shows hybrid reaching the target faster, making this the most robust finding of the benchmark.

Wall-clock time budget

Including the pretraining overhead, the total hybrid wall time to target is:

$$t_{\mathrm{hybrid}}^{\mathrm{total}} = t_{\mathrm{pretrain}} + t_{\mathrm{VMC\ to\ target}} = 2.29 + 4.26 = 6.55\ \mathrm{s}$$

This is comparable to the cold-start time-to-target of 6.52 s. The iteration speedup translates to a roughly break-even wall-clock comparison at $N = 12$, because the pretraining overhead ($\sim$ 2.3 s) nearly offsets the saved VMC iterations. The wall-clock advantage is expected to grow with system size, where VMC iterations become more expensive relative to the fixed-cost NLL pretraining.

---

5. Pretraining Design Choices and Their Impact

Evolution across experimental iterations

The final pipeline configuration was arrived at through systematic iteration. The key design decisions and their effects:

Design choice	Early versions	Final version	Impact
Snapshot count M	1000	1500	Reduced NLL noise, smoother holdout convergence
Pretraining iterations	120–150	60	Holdout NLL plateaus by iter ~ 50; extra iterations add noise without improving generalization
Pretraining LR schedule	Fixed or slow decay	$0.0025 \rightarrow 0.0008$	Cosine-like decay stabilizes late-stage pretraining
Holdout monitoring	NLL proxy only	90/10 train/holdout split	Provides reliable generalization diagnostic; NLL proxy oscillates throughout
Early stopping (VMC)	Aggressive (stopped at 100–150 iters)	Disabled (full 250 iters)	Eliminated premature convergence; improved final accuracy
Hybrid VMC LR	0.005	0.005	Conservative LR preserves pretrained structure while allowing refinement

Design choice	Early versions	Final version	Impact
Cold-start VMC LR	0.008	0.008	Higher LR appropriate for escaping random initialization

The most impactful single change was **cutting pretraining from 150 to 60 iterations**. The holdout NLL curves (Figure 3, panel a) show convergence by iteration 40–50 across all seeds; the remaining iterations in earlier versions added noise to the pretrained parameters without improving generalization. Shorter pretraining yields cleaner initial weights that the VMC refines more effectively, and the pretrained energies (~ -14.6) are actually slightly worse than with longer pretraining (~ -14.9), confirming that the VMC phase does the heavy lifting and only needs a "good enough" initialization.

Pretraining captures amplitudes, VMC corrects phases

The NLL pretraining minimizes $-\sum_{\sigma\in\mathcal{D}} \ln|\psi(\sigma)|^2$, which trains the Born probability distribution (amplitudes squared) but provides no information about the wavefunction phases. For the 1D TFIM with $h > 0$, the ground state is not positive-definite in the computational basis, so the sign structure must be learned during VMC. This division of labor — amplitudes from snapshots, phases from energy minimization — is the conceptual basis of the hybrid approach and explains why the pretrained energy ($\sim 95\%$ of exact) improves rapidly under VMC refinement.

---

6. Discussion

Why snapshot pretraining helps

The core mechanism is dimensionality reduction of the VMC optimization landscape. A randomly initialized RBM starts at $E \approx -12.0$ ($\sim 78\%$ of exact), far from the ground state in parameter space. The VMC+SR optimizer must first navigate a high-curvature region of the energy landscape to reach the basin of attraction around the ground state, consuming ~ 100–150 iterations. The pretrained RBM starts at $E \approx -14.6$ ($\sim 95\%$ of exact), already within the basin, and the VMC+SR only needs to refine the wavefunction — a simpler optimization problem that converges in ~ 80 iterations.

Limitations at $N = 12$

At $N = 12$, the Hilbert space dimension is $2^{12} = 4096$, small enough that cold-start VMC with SR converges reliably within 250 iterations. The pretraining advantage, while consistent and statistically robust (5/5 seeds faster, 4/5 seeds more accurate), is moderate in absolute terms:

a 1.9× iteration speedup and 19% accuracy improvement. The wall-clock speedup is roughly break-even due to the pretraining overhead.

The snapshot generation step currently requires exact diagonalization, which limits the method to system sizes where ED is tractable. For the hybrid approach to be practical beyond ED limits, the snapshots would need to come from alternative sources: quantum gas microscope measurements, tensor network sampling (e.g., DMRG-sampled configurations), or approximate classical simulation methods.

Connection to the NQS portfolio

This benchmark demonstrates the first component of the hybrid classical–quantum pipeline for neural quantum states: using classical simulation data to initialize a variational wavefunction before energy-based optimization. The same pretraining strategy applies to bosonic systems (Bose–Hubbard model with quantum gas microscope snapshots), which is the target application for the FOR5919 MLCQS research program. The key transferable insight is that NLL pretraining on Born-sampled snapshots provides a reliable warm start for VMC, with the convergence speedup expected to grow as system size increases and cold-start initialization becomes progressively worse.

Summary

Deliverable	Corrected summary	Status
Snapshot generation	1500 Born-sampled Z-basis snapshots generated from the ED ground state; reported magnetization is consistent with the symmetric ground state.	□ Validated
NLL pretraining	Across 5 seeds, holdout NLL drops from about 8.3 and plateaus by about iteration 40-50, with final values near 5.2-6.3.	□ Converged
Post-pretrain state	Mean post-pretrain energy is $E_{\text{pretrain}} = -14.614$, or about 95.4% of the exact ground-state energy. Mean pretraining time is 2.29 s for 60 iterations.	□ Validated

Deliverable	Corrected summary	Status
VMC refinement speed	Hybrid reaches $\lvert E - E_{\mathrm{ED}} \rvert \leq 10^{-3}$ in 79.8 iterations on average versus 152.2 for cold start, a 1.91x reduction in VMC iterations. Hybrid is faster in 5/5 seeds on this metric.	□ Confirmed
Final accuracy	Mean final absolute error is 3.62×10^{-3} for hybrid versus 4.47×10^{-3} for cold start (19% lower). Hybrid has lower final error in 4/5 seeds.	□ Confirmed
Threshold crossing	All 5 hybrid seeds cross the 10^{-3} error threshold at some point during optimization; this does not mean the final reported energies are all below 10^{-3}.	□ Confirmed
Wall-clock cost	VMC-only mean time to target is 4.26 s for hybrid versus 6.52 s for cold start, but including the 2.29 s pretraining overhead gives a total hybrid time of about 6.55 s, essentially break-even with cold start at $N = 12$.	□ Clarified
Robustness	Over seeds 42-46, the hybrid method is consistently faster by iteration count and usually, but not always, more accurate at the end of 250 VMC iterations.	□ Robust
ED cross-check	The benchmark uses a fixed ED reference of $E_{\mathrm{ED}} = -15.3226$, consistent across all runs.	□ Exact match

Across five random seeds, snapshot pretraining provides a reliable warm start for the RBM: it substantially reduces the number of VMC+SR iterations needed to reach the target error

and modestly improves the final mean energy error relative to cold-start optimization. At this system size, the clearest benefit is the reduction in optimization iterations rather than a net end-to-end runtime gain, because the pretraining overhead nearly cancels the saved VMC time. The most defensible conclusion is therefore that hybrid snapshot pretraining is a robust initialization strategy for this 1D TFIM benchmark, with stronger practical runtime advantages likely only if the downstream VMC cost grows with system size.